

```
#import file
from google.colab import files
uploaded = files.upload()

Choose Files Corona_NLP_test.csv
• Corona_NLP_test.csv(text/csv) - 1002494 bytes, last modified: 4/21/2023 - 100% done
Saving Corona_NLP_test.csv to Corona_NLP_test.csv

import io
import pandas as pd

#preprocess data
df = pd.read_csv(io.BytesIO(uploaded['Corona_NLP_test.csv']))
df1 = df.drop(['UserName', 'ScreenName', 'Location', 'TweetAt'], axis = 1)
df1.Sentiment.replace(('Extremely Positive','Positive', 'Negative', 'Extremely Negative', 'Neutral'), ('positive', 'positive', 'negative', 'n
df1

OriginalTweet Sentiment
0 TRENDING: New Yorkers encounter empty supermar... negative
1 When I couldn't find hand sanitizer at Fred Me... positive
2 Find out how you can protect yourself and love... positive
3 #Panic buying hits #NewYork City as anxious sh... negative
4 #toiletpaper #dunnypaper #coronavirus #coronav... positive
...
3793 Meanwhile In A Supermarket in Israel -- People... positive
3794 Did you panic buy a lot of non-perishable item... negative
3795 Asst Prof of Economics @cconces was on @NBCPhi... positive
3796 Gov need to do somethings instead of biar je r... negative
3797 I and @ForestandPaper members are committed to... positive
3798 rows x 2 columns


df1.head()

OriginalTweet Sentiment
0 TRENDING: New Yorkers encounter empty supermar... negative
1 When I couldn't find hand sanitizer at Fred Me... positive
2 Find out how you can protect yourself and love... positive
3 #Panic buying hits #NewYork City as anxious sh... negative
4 #toiletpaper #dunnypaper #coronavirus #coronav... positive

#data visualization
import seaborn as sns

sns.countplot(data = df1, x = 'Sentiment')
```

<Axes: xlabel='Sentiment', ylabel='count'>



Here we can see the distribution of the data from the dataset that was chosen. I chose a dataset from kaggle.

<https://www.kaggle.com/datasets/datatattle/covid-19-nlp-text-classification?resource=download>. This dataset looks at COVID related tweets and performs sentiment analysis on them. Here as we can see, there are a larger amount of positive tweets within the dataset but not enough to skew it in any sort of way

```
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras import layers, models
from sklearn.preprocessing import LabelEncoder
import pickle
import numpy as np
import pandas as pd
np.random.seed(1234)

print('rows and columns:', df.shape)

rows and columns: (3798, 6)

#split dataset into train and test samples
i = np.random.rand(len(df)) < 0.8
train = df1[i]
test = df1[~i]
print("train data size: ", train.shape)
print("test data size: ", test.shape)

train data size: (3027, 2)
test data size: (771, 2)

#set up X and Y train and test
num_labels = 2
vocab_size = 25000
batch_size = 100

tokenizer = Tokenizer(num_words=vocab_size)
tokenizer.fit_on_texts(train.OriginalTweet)

x_train = tokenizer.texts_to_matrix(train.OriginalTweet, mode='tfidf')
x_test = tokenizer.texts_to_matrix(test.OriginalTweet, mode='tfidf')

#set up X and Y train and test
encoder = LabelEncoder()
encoder.fit(train.Sentiment)
y_train = encoder.transform(train.Sentiment)
y_test = encoder.transform(test.Sentiment)

print("train shapes:", x_train.shape, y_train.shape)
print("test shapes:", x_test.shape, y_test.shape)
print("test first five labels:", y_test[:5])

train shapes: (3027, 25000) (3027,)
test shapes: (771, 25000) (771,)
test first five labels: [1 0 1 1 0]

# Sequential model fit
model = models.Sequential()
model.add(layers.Dense(32, input_dim=vocab_size, kernel_initializer='normal', activation='relu'))
model.add(layers.Dense(1, kernel_initializer='normal', activation='sigmoid'))
model.compile(loss='binary_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])
history = model.fit(x_train, y_train,
                  batch_size=batch_size,
```

```
epochs=30,
verbose=1,
validation_split=0.1)
```

```
Epoch 2/30
28/28 [=====] - 1s 20ms/step - loss: 0.5718 - accuracy: 0.7419 - val_loss: 0.6143 - val_accuracy: 0.6931
Epoch 3/30
28/28 [=====] - 1s 18ms/step - loss: 0.3922 - accuracy: 0.9214 - val_loss: 0.5698 - val_accuracy: 0.7129
Epoch 4/30
28/28 [=====] - 1s 19ms/step - loss: 0.2181 - accuracy: 0.9666 - val_loss: 0.5619 - val_accuracy: 0.7294
Epoch 5/30
28/28 [=====] - 0s 13ms/step - loss: 0.1163 - accuracy: 0.9890 - val_loss: 0.5718 - val_accuracy: 0.7360
Epoch 6/30
28/28 [=====] - 0s 12ms/step - loss: 0.0668 - accuracy: 0.9971 - val_loss: 0.5828 - val_accuracy: 0.7360
Epoch 7/30
28/28 [=====] - 0s 13ms/step - loss: 0.0421 - accuracy: 0.9993 - val_loss: 0.5964 - val_accuracy: 0.7327
Epoch 8/30
28/28 [=====] - 0s 12ms/step - loss: 0.0286 - accuracy: 0.9996 - val_loss: 0.6127 - val_accuracy: 0.7393
Epoch 9/30
28/28 [=====] - 0s 12ms/step - loss: 0.0206 - accuracy: 1.0000 - val_loss: 0.6292 - val_accuracy: 0.7393
Epoch 10/30
28/28 [=====] - 0s 12ms/step - loss: 0.0156 - accuracy: 1.0000 - val_loss: 0.6455 - val_accuracy: 0.7393
Epoch 11/30
28/28 [=====] - 0s 12ms/step - loss: 0.0121 - accuracy: 1.0000 - val_loss: 0.6589 - val_accuracy: 0.7360
Epoch 12/30
28/28 [=====] - 0s 12ms/step - loss: 0.0097 - accuracy: 1.0000 - val_loss: 0.6733 - val_accuracy: 0.7393
Epoch 13/30
28/28 [=====] - 0s 13ms/step - loss: 0.0079 - accuracy: 1.0000 - val_loss: 0.6862 - val_accuracy: 0.7360
Epoch 14/30
28/28 [=====] - 0s 12ms/step - loss: 0.0065 - accuracy: 1.0000 - val_loss: 0.6979 - val_accuracy: 0.7360
Epoch 15/30
28/28 [=====] - 0s 12ms/step - loss: 0.0055 - accuracy: 1.0000 - val_loss: 0.7089 - val_accuracy: 0.7327
Epoch 16/30
28/28 [=====] - 0s 12ms/step - loss: 0.0047 - accuracy: 1.0000 - val_loss: 0.7195 - val_accuracy: 0.7360
Epoch 17/30
28/28 [=====] - 0s 12ms/step - loss: 0.0041 - accuracy: 1.0000 - val_loss: 0.7290 - val_accuracy: 0.7294
Epoch 18/30
28/28 [=====] - 0s 12ms/step - loss: 0.0036 - accuracy: 1.0000 - val_loss: 0.7393 - val_accuracy: 0.7327
Epoch 19/30
28/28 [=====] - 0s 13ms/step - loss: 0.0031 - accuracy: 1.0000 - val_loss: 0.7497 - val_accuracy: 0.7327
Epoch 20/30
28/28 [=====] - 0s 12ms/step - loss: 0.0028 - accuracy: 1.0000 - val_loss: 0.7590 - val_accuracy: 0.7327
Epoch 21/30
28/28 [=====] - 0s 12ms/step - loss: 0.0025 - accuracy: 1.0000 - val_loss: 0.7670 - val_accuracy: 0.7360
Epoch 22/30
28/28 [=====] - 0s 13ms/step - loss: 0.0022 - accuracy: 1.0000 - val_loss: 0.7755 - val_accuracy: 0.7327
Epoch 23/30
28/28 [=====] - 0s 12ms/step - loss: 0.0020 - accuracy: 1.0000 - val_loss: 0.7831 - val_accuracy: 0.7327
Epoch 24/30
28/28 [=====] - 0s 12ms/step - loss: 0.0018 - accuracy: 1.0000 - val_loss: 0.7909 - val_accuracy: 0.7360
Epoch 25/30
28/28 [=====] - 0s 12ms/step - loss: 0.0017 - accuracy: 1.0000 - val_loss: 0.7983 - val_accuracy: 0.7327
Epoch 26/30
28/28 [=====] - 0s 12ms/step - loss: 0.0015 - accuracy: 1.0000 - val_loss: 0.8057 - val_accuracy: 0.7327
Epoch 27/30
28/28 [=====] - 0s 12ms/step - loss: 0.0014 - accuracy: 1.0000 - val_loss: 0.8124 - val_accuracy: 0.7327
Epoch 28/30
28/28 [=====] - 0s 12ms/step - loss: 0.0013 - accuracy: 1.0000 - val_loss: 0.8186 - val_accuracy: 0.7327
Epoch 29/30
28/28 [=====] - 0s 12ms/step - loss: 0.0012 - accuracy: 1.0000 - val_loss: 0.8255 - val_accuracy: 0.7327
Epoch 30/30
28/28 [=====] - 0s 12ms/step - loss: 0.0011 - accuracy: 1.0000 - val_loss: 0.8313 - val_accuracy: 0.7327
```

```
#evaluation
```

```
score = model.evaluate(x_test, y_test, batch_size=batch_size, verbose=1)
print('Accuracy: ', score[1])
```

```
8/8 [=====] - 0s 4ms/step - loss: 0.9157 - accuracy: 0.6952
Accuracy: 0.69520103931427
```

```
print(score)
```

```
[0.9156782031059265, 0.69520103931427]
```

```
pred = model.predict(x_test)
pred_labels = [1 if p>0.5 else 0 for p in pred]
```

```
25/25 [=====] - 0s 2ms/step
```

```

from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
print('accuracy score: ', accuracy_score(y_test, pred_labels))
print('precision score: ', precision_score(y_test, pred_labels))
print('recall score: ', recall_score(y_test, pred_labels))
print('f1 score: ', f1_score(y_test, pred_labels))

accuracy score: 0.695201037613489
precision score: 0.706140350877193
recall score: 0.7612293144208038
f1 score: 0.732650739476678

```

As we can see here, using a sequential model in terms of this selected dataset, we have reached an accuracy score of 0.6952. A precision of 0.706, recall of 0.76123, and f1 score of 0.7327. As a result, we can see that this model is pretty accurate as a whole.

```

from tensorflow.keras import layers, models, preprocessing

max_features = 10000
maxlen = 500
batch_size = 32

#split dataset into train and test
i = np.random.rand(len(df)) < 0.8
train = df1[i]
test = df1[~i]
print("train data size: ", train.shape)
print("test data size: ", test.shape)

train data size: (3040, 2)
test data size: (758, 2)

#set up X and Y train and test
tokenizer = Tokenizer(num_words=vocab_size)
tokenizer.fit_on_texts(train.OriginalTweet)

x_train = tokenizer.texts_to_matrix(train.OriginalTweet, mode='tfidf')
x_test = tokenizer.texts_to_matrix(test.OriginalTweet, mode='tfidf')

encoder = LabelEncoder()
encoder.fit(train.Sentiment)
y_train = encoder.transform(train.Sentiment)
y_test = encoder.transform(test.Sentiment)

train_data = preprocessing.sequence.pad_sequences(x_train, maxlen=maxlen)
test_data = preprocessing.sequence.pad_sequences(x_test, maxlen=maxlen)

train_data.shape

(3040, 500)

#build CNN network
model = models.Sequential()
model.add(layers.Embedding(max_features, 128, input_length=maxlen))
model.add(layers.Conv1D(32, 7, activation='relu'))
model.add(layers.MaxPooling1D(5))
model.add(layers.Conv1D(32, 7, activation='relu'))
model.add(layers.GlobalMaxPooling1D())
model.add(layers.Dense(1))

model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 500, 128)	1280000
conv1d (Conv1D)	(None, 494, 32)	28704
max_pooling1d (MaxPooling1D)	(None, 98, 32)	0
conv1d_1 (Conv1D)	(None, 92, 32)	7200
global_max_pooling1d (GlobalMaxPooling1D)	(None, 32)	0

```

lMaxPooling1D)

dense_2 (Dense)          (None, 1)          33

=====
Total params: 1,315,937
Trainable params: 1,315,937
Non-trainable params: 0
=====

model.compile(optimizer=tf.keras.optimizers.RMSprop(learning_rate=1e-4), # set learning rate
              loss='binary_crossentropy',
              metrics=['accuracy'])

history = model.fit(train_data,
                    y_train,
                    epochs=10,
                    batch_size=128,
                    validation_split=0.2)

Epoch 1/10
19/19 [=====] - 7s 349ms/step - loss: 1.4804 - accuracy: 0.4408 - val_loss: 1.2235 - val_accuracy: 0.3964
Epoch 2/10
19/19 [=====] - 6s 296ms/step - loss: 0.9981 - accuracy: 0.4408 - val_loss: 0.9064 - val_accuracy: 0.3964
Epoch 3/10
19/19 [=====] - 6s 339ms/step - loss: 0.7816 - accuracy: 0.4408 - val_loss: 0.7249 - val_accuracy: 0.3964
Epoch 4/10
19/19 [=====] - 6s 297ms/step - loss: 0.6949 - accuracy: 0.5090 - val_loss: 0.6761 - val_accuracy: 0.6036
Epoch 5/10
19/19 [=====] - 6s 336ms/step - loss: 0.6866 - accuracy: 0.5592 - val_loss: 0.6740 - val_accuracy: 0.6036
Epoch 6/10
19/19 [=====] - 6s 297ms/step - loss: 0.6864 - accuracy: 0.5592 - val_loss: 0.6865 - val_accuracy: 0.6036
Epoch 7/10
19/19 [=====] - 6s 337ms/step - loss: 0.6875 - accuracy: 0.5592 - val_loss: 0.6733 - val_accuracy: 0.6036
Epoch 8/10
19/19 [=====] - 6s 294ms/step - loss: 0.6867 - accuracy: 0.5592 - val_loss: 0.6742 - val_accuracy: 0.6036
Epoch 9/10
19/19 [=====] - 6s 338ms/step - loss: 0.6874 - accuracy: 0.5592 - val_loss: 0.6750 - val_accuracy: 0.6036
Epoch 10/10
19/19 [=====] - 6s 295ms/step - loss: 0.6874 - accuracy: 0.5592 - val_loss: 0.6779 - val_accuracy: 0.6036

from sklearn import metrics

pred = model.predict(test_data)
pred = [1.0 if p>= 0.01 else 0 for p in pred]

24/24 [=====] - 0s 2ms/step

print(metrics.f1_score(y_test, pred, average='weighted'))
print(metrics.precision_score(y_test, pred, average='weighted'))
print(metrics.recall_score(y_test, pred, average='weighted'))

0.3936873932511656
0.3056742812391915
0.5528781793842035
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision is ill-defined and bei
_warn_prf(average, modifier, msg_start, len(result))

```

Here, using the CNN architecture, we get an f1 score of 0.3937, precision score of 0.3057, and recall score 0.5529

```

#split dataset into train and test
max_features = 10000
maxlen = 500
batch_size = 32

i = np.random.rand(len(df)) < 0.8
train = df1[i]
test = df1[~i]
print("train data size: ", train.shape)
print("test data size: ", test.shape)

train data size: (3056, 2)
test data size: (742, 2)

```

```
#set up X and Y train and test
tokenizer = Tokenizer(num_words=vocab_size)
tokenizer.fit_on_texts(train.OriginalTweet)

x_train = tokenizer.texts_to_matrix(train.OriginalTweet, mode='tfidf')
x_test = tokenizer.texts_to_matrix(test.OriginalTweet, mode='tfidf')

encoder = LabelEncoder()
encoder.fit(train.Sentiment)
y_train = encoder.transform(train.Sentiment)
y_test = encoder.transform(test.Sentiment)

train_data = preprocessing.sequence.pad_sequences(x_train, maxlen=maxlen)
test_data = preprocessing.sequence.pad_sequences(x_test, maxlen=maxlen)
```

```
#build RNN network architecture
model = models.Sequential()
model.add(layers.Embedding(max_features, 32))
model.add(layers.SimpleRNN(32))
model.add(layers.Dense(1, activation='sigmoid'))
```

```
model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
embedding_1 (Embedding)	(None, None, 32)	320000
simple_rnn (SimpleRNN)	(None, 32)	2080
dense_3 (Dense)	(None, 1)	33
=====		
Total params: 322,113		
Trainable params: 322,113		
Non-trainable params: 0		

```
model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['accuracy'])
```

```
history = model.fit(train_data,
                    y_train,
                    epochs=10,
                    batch_size=128,
                    validation_split=0.2)
```

```
Epoch 1/10
20/20 [=====] - 3s 98ms/step - loss: 0.6866 - accuracy: 0.5615 - val_loss: 0.6686 - val_accuracy: 0.6105
Epoch 2/10
20/20 [=====] - 2s 89ms/step - loss: 0.6856 - accuracy: 0.5656 - val_loss: 0.6691 - val_accuracy: 0.6105
Epoch 3/10
20/20 [=====] - 2s 89ms/step - loss: 0.6852 - accuracy: 0.5656 - val_loss: 0.6712 - val_accuracy: 0.6105
Epoch 4/10
20/20 [=====] - 2s 89ms/step - loss: 0.6855 - accuracy: 0.5656 - val_loss: 0.6760 - val_accuracy: 0.6105
Epoch 5/10
20/20 [=====] - 2s 88ms/step - loss: 0.6848 - accuracy: 0.5656 - val_loss: 0.6795 - val_accuracy: 0.6105
Epoch 6/10
20/20 [=====] - 3s 132ms/step - loss: 0.6857 - accuracy: 0.5656 - val_loss: 0.6704 - val_accuracy: 0.6105
Epoch 7/10
20/20 [=====] - 2s 87ms/step - loss: 0.6849 - accuracy: 0.5656 - val_loss: 0.6739 - val_accuracy: 0.6105
Epoch 8/10
20/20 [=====] - 2s 88ms/step - loss: 0.6855 - accuracy: 0.5656 - val_loss: 0.6707 - val_accuracy: 0.6105
Epoch 9/10
20/20 [=====] - 2s 89ms/step - loss: 0.6851 - accuracy: 0.5656 - val_loss: 0.6703 - val_accuracy: 0.6105
Epoch 10/10
20/20 [=====] - 2s 90ms/step - loss: 0.6853 - accuracy: 0.5656 - val_loss: 0.6706 - val_accuracy: 0.6105
```

```
print('f1 Score: ', metrics.f1_score(y_test, pred, average='weighted'))
print('Precision Score: ', metrics.precision_score(y_test, pred, average='weighted'))
print('Recall Score: ', metrics.recall_score(y_test, pred, average='weighted'))
```

```
f1 Score: 0.3936873932511656
Precision Score: 0.3056742812391915
Recall Score: 0.5528781793842035
```

```
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision is ill-defined and bei
_warn_prf(average, modifier, msg_start, len(result))
```

Using a RNN network, we have an f1 score of 0.3937, precision score of 0.3057, and recall score of 0.5529

```
max_features = 10000
maxlen = 20
batch_size = 32

i = np.random.rand(len(df)) < 0.8
train = df1[i]
test = df1[~i]
print("train data size: ", train.shape)
print("test data size: ", test.shape)

train data size: (3051, 2)
test data size: (747, 2)

tokenizer = Tokenizer(num_words=vocab_size)
tokenizer.fit_on_texts(train.OriginalTweet)

x_train = tokenizer.texts_to_matrix(train.OriginalTweet, mode='tfidf')
x_test = tokenizer.texts_to_matrix(test.OriginalTweet, mode='tfidf')

encoder = LabelEncoder()
encoder.fit(train.Sentiment)
y_train = encoder.transform(train.Sentiment)
y_test = encoder.transform(test.Sentiment)

train_data = preprocessing.sequence.pad_sequences(x_train, maxlen=maxlen)
test_data = preprocessing.sequence.pad_sequences(x_test, maxlen=maxlen)

model = models.Sequential()
model.add(layers.Embedding(max_features, 8, input_length=maxlen))
model.add(layers.Flatten())
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.summary()

history = model.fit(train_data, y_train, epochs=10, batch_size=32, validation_split=0.2)

Model: "sequential_8"

```

Layer (type)	Output Shape	Param #
embedding_7 (Embedding)	(None, 20, 8)	80000
flatten_4 (Flatten)	(None, 160)	0
dense_11 (Dense)	(None, 16)	2576
dense_12 (Dense)	(None, 1)	17

```

Total params: 82,593
Trainable params: 82,593
Non-trainable params: 0

Epoch 1/10
77/77 [=====] - 1s 3ms/step - loss: 0.6876 - acc: 0.5545 - val_loss: 0.6704 - val_acc: 0.6236
Epoch 2/10
77/77 [=====] - 0s 2ms/step - loss: 0.6859 - acc: 0.5619 - val_loss: 0.6712 - val_acc: 0.6236
Epoch 3/10
77/77 [=====] - 0s 2ms/step - loss: 0.6858 - acc: 0.5619 - val_loss: 0.6701 - val_acc: 0.6236
Epoch 4/10
77/77 [=====] - 0s 2ms/step - loss: 0.6855 - acc: 0.5619 - val_loss: 0.6666 - val_acc: 0.6236
Epoch 5/10
77/77 [=====] - 0s 2ms/step - loss: 0.6859 - acc: 0.5619 - val_loss: 0.6725 - val_acc: 0.6236
Epoch 6/10
77/77 [=====] - 0s 2ms/step - loss: 0.6859 - acc: 0.5619 - val_loss: 0.6701 - val_acc: 0.6236
Epoch 7/10
77/77 [=====] - 0s 3ms/step - loss: 0.6858 - acc: 0.5619 - val_loss: 0.6708 - val_acc: 0.6236
Epoch 8/10
77/77 [=====] - 0s 2ms/step - loss: 0.6858 - acc: 0.5619 - val_loss: 0.6681 - val_acc: 0.6236

```

```
Epoch 9/10
77/77 [=====] - 0s 2ms/step - loss: 0.6859 - acc: 0.5619 - val_loss: 0.6720 - val_acc: 0.6236
Epoch 10/10
77/77 [=====] - 0s 2ms/step - loss: 0.6859 - acc: 0.5619 - val_loss: 0.6730 - val_acc: 0.6236
```

```
print('f1 Score: ', metrics.f1_score(y_test, pred, average='weighted'))
print('Precision Score: ', metrics.precision_score(y_test, pred, average='weighted'))
print('Recall Score: ', metrics.recall_score(y_test, pred, average='weighted'))

f1 Score: 0.3936873932511656
Precision Score: 0.3056742812391915
Recall Score: 0.5528781793842035
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision is ill-defined and bei
_warn_prf(average, modifier, msg_start, len(result))
```

Using an embedded approach, we have an f1 score of 0.3937, precision score of 0.3057, and recall score of 0.5529